

SUPERVISED MACHINE LEARNING

LINEAR REGRESSION

LINEAR REGRESSION

Step	Process
1	Define model ($y = mx + b$)
2	Compute loss (MSE)
3	Compute gradients of loss
4	Update parameters
5	Repeat for many epochs

LINE BY LINE CODE EXECUTION

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

- **pandas** → Used for handling datasets (reading CSV, manipulating tables, etc.).
- **numpy** → Used for numerical operations (arrays, linear algebra, etc.).
- **matplotlib.pyplot** → Used for plotting graphs and visualizations.

READ DATA

- `data = pd.read_csv('studytime_scores.csv')`
- Reads the CSV file `studytime_scores.csv` into a pandas **DataFrame** called `data`.
- The dataset contains two main columns: `StudyTime_Hours` → Number of hours a student studied (< 9000 hours).
- `ExamScore` → Score obtained (< 300).

DATA NORMALIZATION

```
data['StudyTime_Hours_Norm'] = data['StudyTime_Hours'] /  
                                data['StudyTime_Hours'].max()
```

- Creates a new column **StudyTime_Hours_Norm** by dividing each value by the **maximum study hours**.
- This normalizes the data into the range [0,1], which helps gradient descent **converge faster**.

LOSS FUNCTION

```
def loss_function(m, b, points):  
    total_error = 0  
  
    for i in range(len(points)):  
        x = points.iloc[i].StudyTime_Hours_Norm  
        y = points.iloc[i].ExamScore  
        total_error += (y - (m*x + b))**2  
  
    return total_error / float(len(points))
```

This function calculates **Mean Squared Error (MSE)**. $m \rightarrow$ slope, $b \rightarrow$ intercept, $\text{points} \rightarrow$ dataset.

For each data point: $y - (m*x + b) \rightarrow$ difference between actual and predicted value.

Square it $(y - (m*x + b))^2$

Sum all squared errors, divide by number of points \rightarrow average error.



```
def gradient_descent(m_now, b_now, points, L):

    m_gradient = 0

    b_gradient = 0

    n = len(points)

    for i in range(n):

        x = points.iloc[i].StudyTime_Hours_Norm

        y = points.iloc[i].ExamScore

        m_gradient += -(2/n) * x * (y - (m_now*x + b_now))

        b_gradient += -(2/n) * (y - (m_now*x + b_now))

    m = m_now - L * m_gradient

    b = b_now - L * b_gradient

    return m, b
```

GRADIENT DESCENT

Computes gradients of the loss function with respect to **m and b**.

$-(2/n) * x * (y - (m*x + b)) \rightarrow$ partial derivative of loss w.r.t m.

$-(2/n) * (y - (m*x + b)) \rightarrow$ partial derivative w.r.t b.

L \rightarrow learning rate; controls how big the step is in each update.

Updates parameters: =

$m = m_now - L * m_gradient$

$b = b_now - L * b_gradient$

Returns the **updated m and b**.

INITIALIZING PARAMETERS

$m = 0$

Start with slope $m=0$ and intercept $b=0$.

$b = 0$

Learning rate $L=0.01 \rightarrow$ small enough for stable convergence.

$L = 0.01$

$\text{epochs} = 2000$

$\text{epochs}=2000 \rightarrow$ number of iterations gradient descent will run.

TRAINING THE MODEL

```
loss_history = []
```

```
for i in range(epochs+1):
```

```
    m, b = gradient_descent(m, b, data, L)
```

```
    current_loss = loss_function(m, b, data)
```

```
    loss_history.append(current_loss)
```

```
    if i % 200 == 0:
```

```
        print(f'Epoch {i}: m={m:.6f}, b={b:.3f},  
loss={loss_function(m,b,data):.3f}')
```

Loop runs gradient descent 2000 times.

Each iteration: Update m and b using gradients.

Compute current loss and store in loss_history.

Every 200 epochs, print current slope, intercept, and loss.



FINAL MODEL

```
print(f"\nFinal model: y = {m:.6f}x  
+ {b:.3f}")
```

- Prints the **final linear regression equation** after training.



PLOTTING RESULTS

```
plt.figure(figsize=(12,5))
```

```
plt.subplot(1, 2, 1)
```

```
plt.scatter(data.StudyTime_Hours_Norm, data.ExamScore,  
color="black", label="Data points")
```

```
x_range = np.linspace(data.StudyTime_Hours_Norm.min(),  
data.StudyTime_Hours_Norm.max(), 100)
```

```
plt.plot(x_range, m*x_range + b, color="red",  
label="Regression line")
```

```
plt.xlabel("Normalized Study Time")
```

```
plt.ylabel("Exam Score")
```

```
plt.title("Linear Regression Fit")
```

```
plt.legend()
```

```
plt.subplot(1, 2, 2)
```

```
plt.plot(range(epochs+1), loss_history,  
color='blue')
```

```
plt.xlabel("Epochs")
```

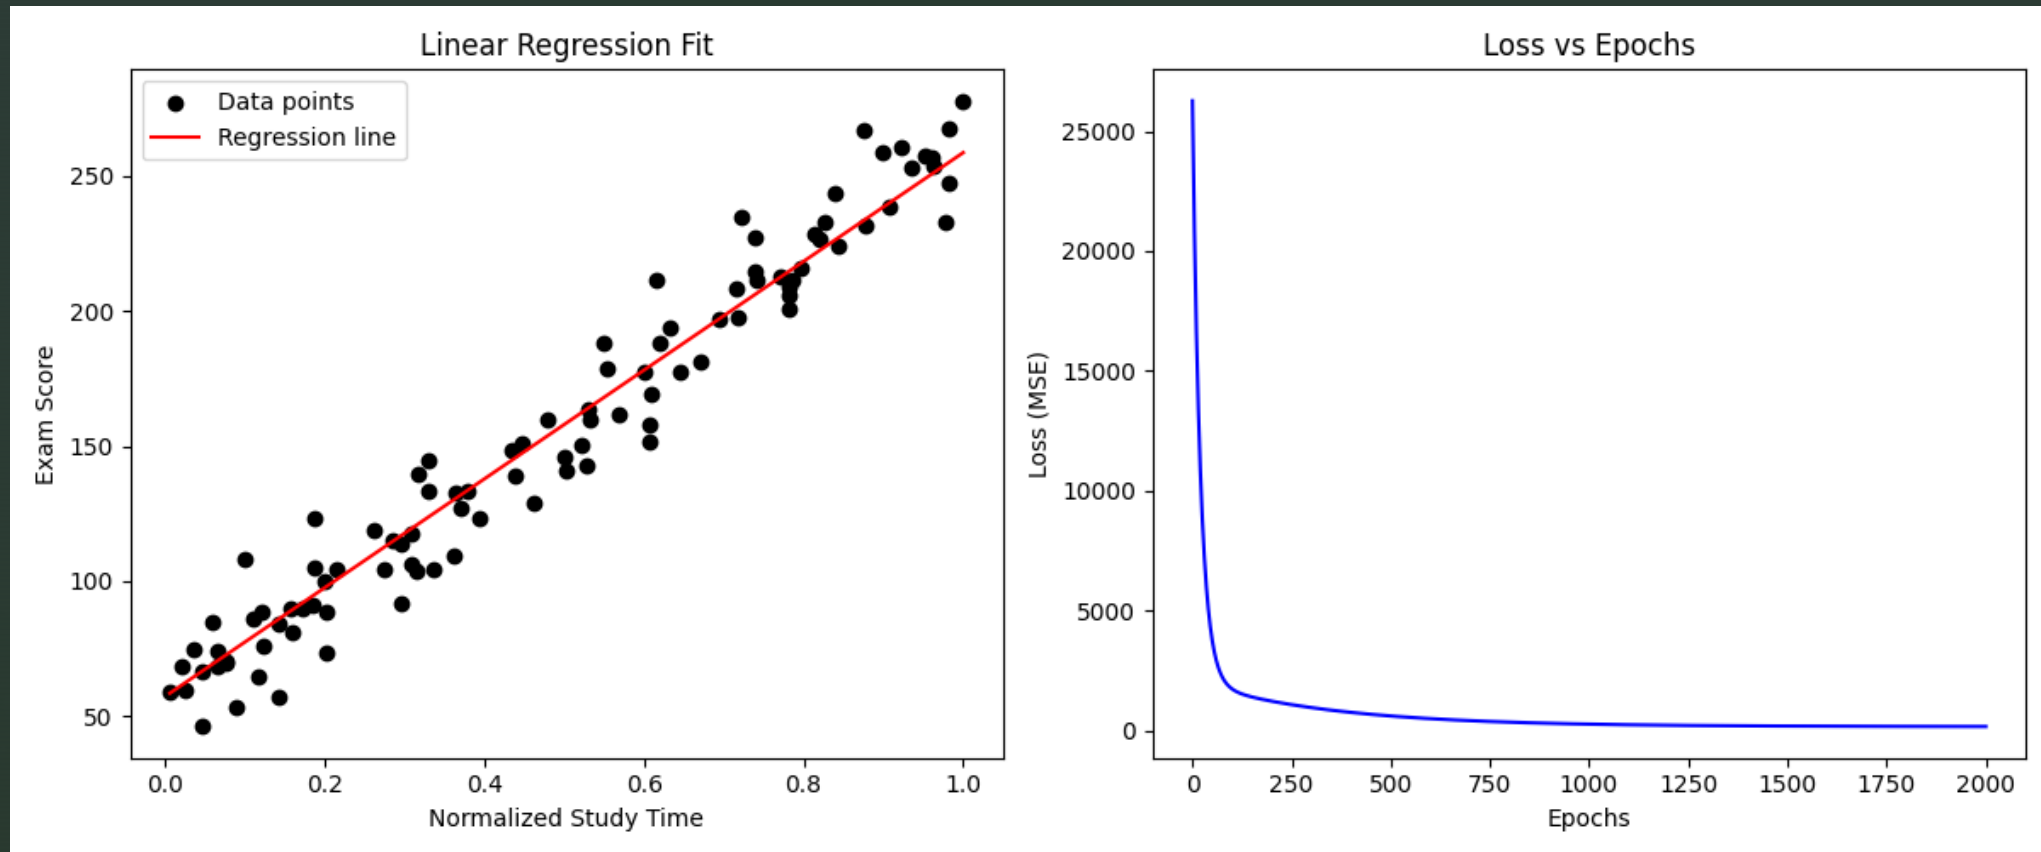
```
plt.ylabel("Loss (MSE)")
```

```
plt.title("Loss vs Epochs")
```

```
plt.tight_layout()
```

```
plt.show()
```

OUTPUT



HOW MODEL GOT TRAINED ?

Epoch 0: $m=1.834035$, $b=3.059$, $\text{loss}=26261.503$

Epoch 200: $m=101.235673$, $b=107.721$, $\text{loss}=1241.455$

Epoch 400: $m=128.694817$, $b=94.626$, $\text{loss}=775.218$

Epoch 600: $m=148.952100$, $b=84.229$, $\text{loss}=514.378$

Epoch 800: $m=164.118378$, $b=76.441$, $\text{loss}=368.133$

Epoch 1000: $m=175.474566$, $b=70.609$, $\text{loss}=286.138$

Epoch 1200: $m=183.977848$, $b=66.242$, $\text{loss}=240.165$

Epoch 1400: $m=190.344932$, $b=62.972$, $\text{loss}=214.390$

Epoch 1600: $m=195.112475$, $b=60.524$, $\text{loss}=199.938$

Epoch 1800: $m=198.682314$, $b=58.691$, $\text{loss}=191.836$

Epoch 2000: $m=201.355337$, $b=57.318$, $\text{loss}=187.293$

Final model: $y = 201.355337x + 57.318$